



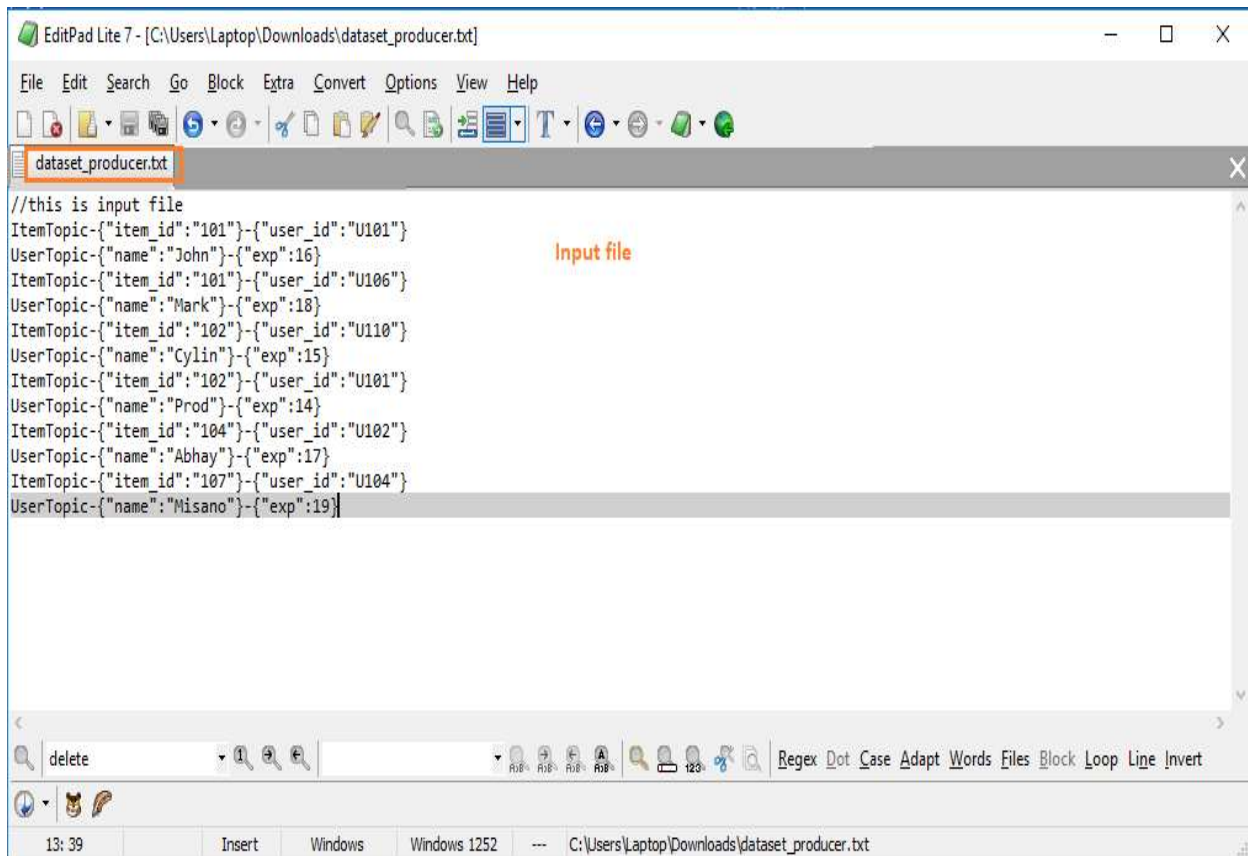
BIG DATA HADOOP & SPARK TRAINING

Assignment on Apache Kafka II

Contents

Input dataset for both the tasks	2
Task 1: Create a java program MyKafkaProducer.java that takes a file name and delimiter as input arguments. It should read the content of file line by line. Fields in the file are in following order	3
1. Kafka Topic Name	3
2. Key	3
3. value	3
For every line, insert the key and value to the respective Kafka broker in a fire and forget mode.	3
After record is sent, it should print appropriate message on screen.	3
Pass dataset_producer.txt as the input file and -as delimiter.	3
LINK:	
https://drive.google.com/file/d/0B_Qjau8wv1KoSnR5eHpKOF9rTFU/view?usp=sharing	3
Task 2: Modify the previous program MyKafkaProducer.java and create a new Java program KafkaProducerWithAck.java. This should perform the same task as of KafkaProducer.java with some modification. When passing any data to a topic, it should wait for acknowledgement. After acknowledgement is received from the broker, it should print the key and value which has been written to a specified topic. The application should attempt for 3 retries before giving any exception.....	8
Pass dataset_producer.txt as the input file and -as delimiter.	8
Program to perform this task is as below:	8

Input dataset for both the tasks



```
//this is input file
ItemTopic-{"item_id":"101"}-{"user_id":"U101"}
UserTopic-{"name":"John"}-{"exp":16}
ItemTopic-{"item_id":"101"}-{"user_id":"U106"}
UserTopic-{"name":"Mark"}-{"exp":18}
ItemTopic-{"item_id":"102"}-{"user_id":"U110"}
UserTopic-{"name":"Cylin"}-{"exp":15}
ItemTopic-{"item_id":"102"}-{"user_id":"U101"}
UserTopic-{"name":"Prod"}-{"exp":14}
ItemTopic-{"item_id":"104"}-{"user_id":"U102"}
UserTopic-{"name":"Abhay"}-{"exp":17}
ItemTopic-{"item_id":"107"}-{"user_id":"U104"}
UserTopic-{"name":"Misano"}-{"exp":19}
```

This file has two topics namely:

1. ItemTopic- It has item_id and user_id
2. UserTopic- It has user details like name and experience

Task 1: Create a java program MyKafkaProducer.java that takes a file name and delimiter as input arguments. It should read the content of file line by line. Fields in the file are in following order

1. Kafka Topic Name

2. Key

3. value

For every line, insert the key and value to the respective Kafka broker in a fire and forget mode.

After record is sent, it should print appropriate message on screen.

Pass dataset_producer.txt as the input file and -as delimiter.

LINK:

https://drive.google.com/file/d/0B_Qjau8wv1KoSnR5eHpKOF9rTFU/view?usp=sharing

Program to perform this task is as below:

- Packages required to be imported are as follows:

```
import org.apache.kafka.clients.producer.KafkaProducer;  
import org.apache.kafka.clients.producer.ProducerRecord;  
import java.io.BufferedReader;  
import java.io.FileReader;  
import java.io.IOException;  
import java.util.Properties;
```

- This is the class called “MyKafkaProducer” which takes two arguments (Input File name and delimiter) in the command line.

```
public class MyKafkaProducer {
    public static void main(String[] args) throws IOException{
        if (args.length != 2) {
            System.out.println("Please provide appropriate command line arguments");
            System.exit(-1);
        }
    }
}
```

- We configure the properties for KafkaProducer:
 - We create a new instance of Properties called *props*.
 - Using this instance we add properties to kafkaProducer like, bootstrapserver/meta-data-brokerlist, key and value serializers.

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
```

- We then instantiate the KafkaProducer class called producer, we have mentioned string in <> because both key and value are String.
- We add the properties instance (props)to KafkaProducer instance.
- We also instantiate ProducerRecord as producerRecord

```
KafkaProducer<String, String> producer = new KafkaProducer<>(props);
ProducerRecord<String, String> producerRecord = null;
```

- Now we take the data provided in the command line i.e. file name and delimiter and save them in the array of string variables called filename and delimiter

```
String fileName = args[0];
String delimiter = args[1];
```

- We read the contents of the input file, and save their contents arrays in different variables:
 - We save the topic name i.e. first part of array(0th index elements) in String variable topic and similarly we save key and value variables too.

```

try(BufferedReader br = new BufferedReader(new FileReader(fileName))) {
    for(String line; (line = br.readLine()) != null; ) {
        String[] tempArray = line.split(delimiter);
        String topic = tempArray[0];
        String key = tempArray[1];
        String value = tempArray[2];

```

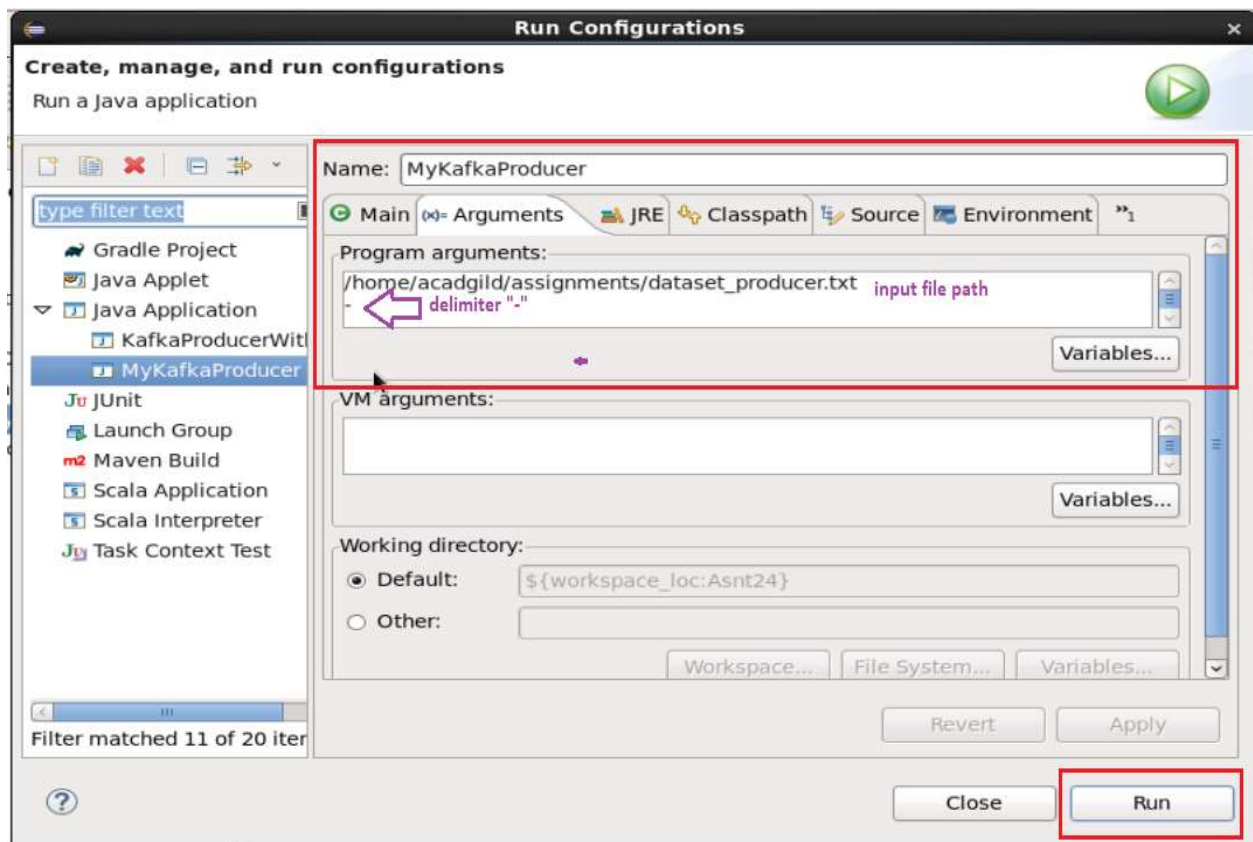
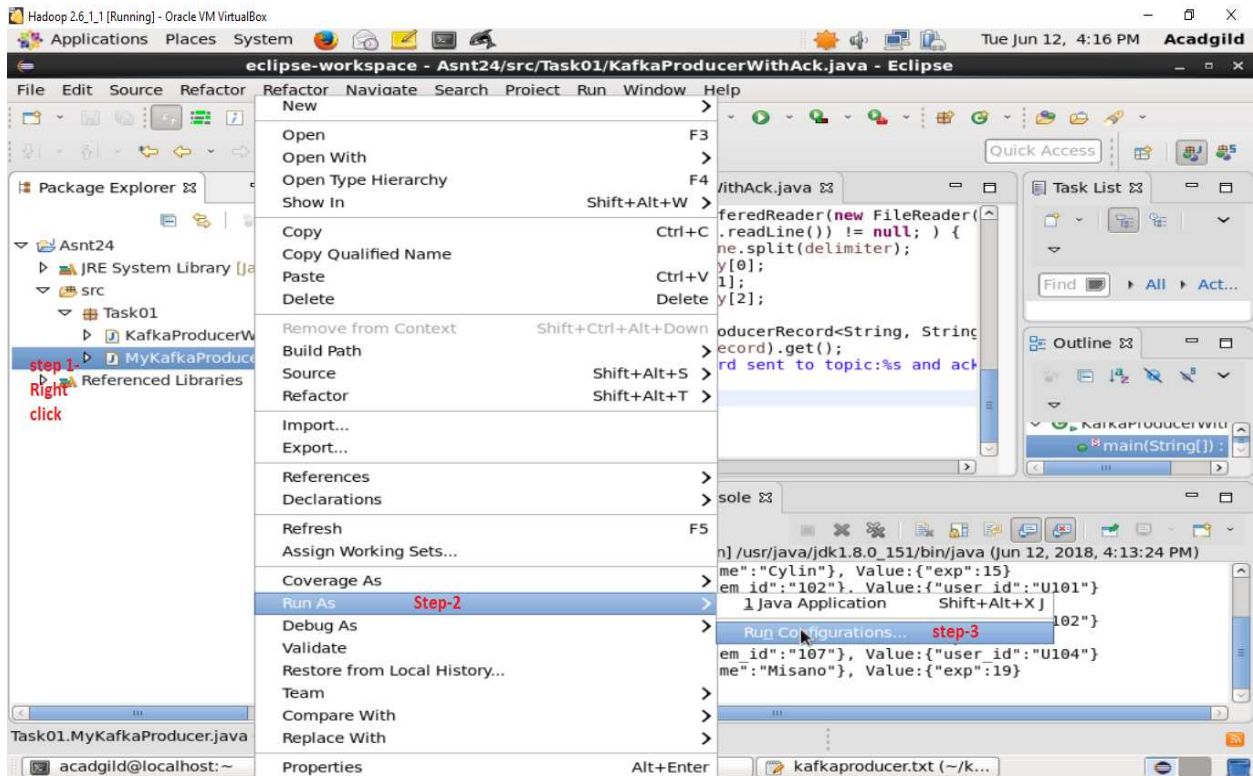
- Now, we pass the variables topic, key and value to producer record.
- We also print appropriate message which shows the topics, key and value contents.
- We finally, close the producer.

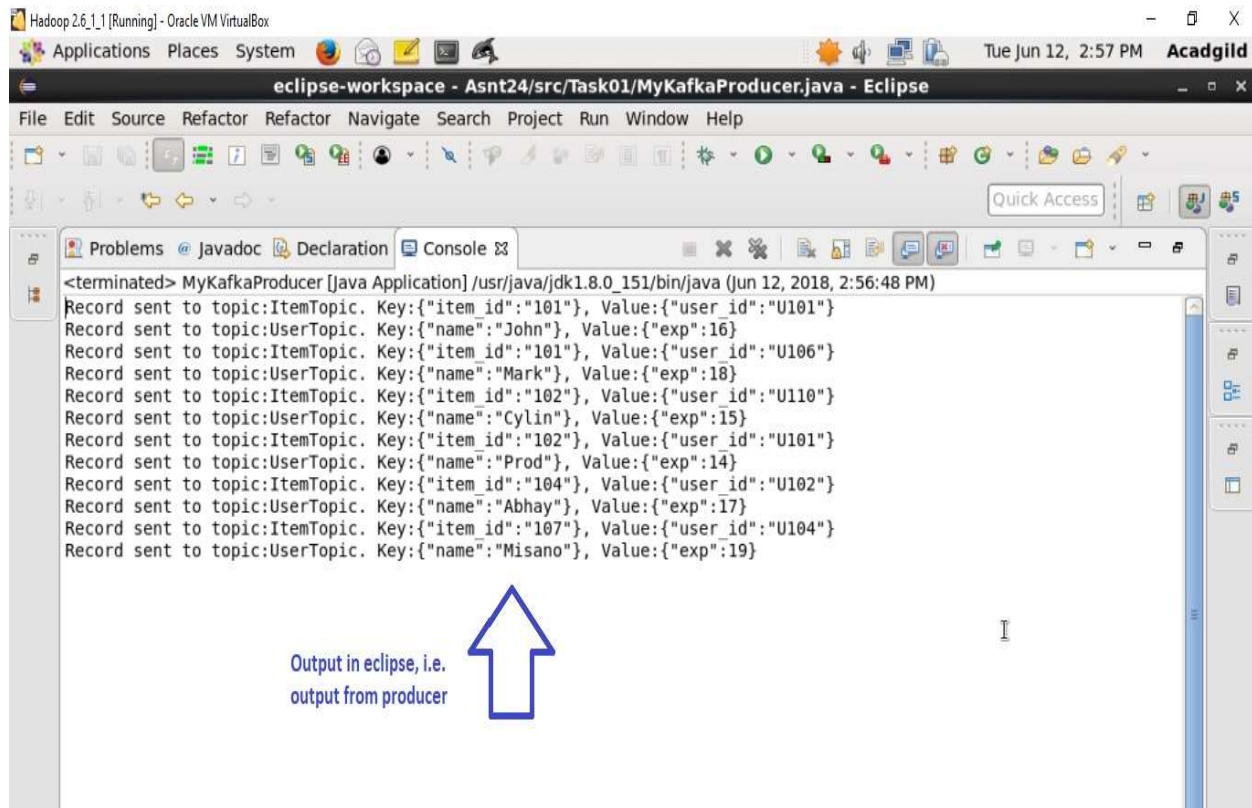
```

        producerRecord = new ProducerRecord<String, String>(topic, key, value);
        producer.send(producerRecord);
        System.out.printf("Record sent to topic:%s. Key:%s, Value:%s\n", topic, key, value);
    } } producer.close(); } }

```

- **We start zookeeper and kafka server using the below commands:**
 - ## starting zookeeper
bin/zookeeper-server-start.sh config/zookeeper.properties
 - ## starting kafka server
bin/kafka-server-start.sh config/server.properties
- Now, we run this program in eclipse, by giving the arguments in “Run Configurations” as shown below:





Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox

Applications Places System Tue Jun 12, 2:57 PM Acadgild

eclipse-workspace - Asnt24/src/Task01/MyKafkaProducer.java - Eclipse

File Edit Source Refactor Refactor Navigate Search Project Run Window Help

Problems Javadoc Declaration Console

```
<terminated> MyKafkaProducer [Java Application] /usr/java/jdk1.8.0_151/bin/java (Jun 12, 2018, 2:56:48 PM)
Record sent to topic:ItemTopic. Key:{"item_id":"101"}, Value:{"user_id":"U101"}
Record sent to topic:UserTopic. Key:{"name":"John"}, Value:{"exp":16}
Record sent to topic:ItemTopic. Key:{"item_id":"101"}, Value:{"user_id":"U106"}
Record sent to topic:UserTopic. Key:{"name":"Mark"}, Value:{"exp":18}
Record sent to topic:ItemTopic. Key:{"item_id":"102"}, Value:{"user_id":"U110"}
Record sent to topic:UserTopic. Key:{"name":"Cylin"}, Value:{"exp":15}
Record sent to topic:ItemTopic. Key:{"item_id":"102"}, Value:{"user_id":"U101"}
Record sent to topic:UserTopic. Key:{"name":"Prod"}, Value:{"exp":14}
Record sent to topic:ItemTopic. Key:{"item_id":"104"}, Value:{"user_id":"U102"}
Record sent to topic:UserTopic. Key:{"name":"Abhay"}, Value:{"exp":17}
Record sent to topic:ItemTopic. Key:{"item_id":"107"}, Value:{"user_id":"U104"}
Record sent to topic:UserTopic. Key:{"name":"Misano"}, Value:{"exp":19}
```

Output in eclipse, i.e. output from producer

- Now, we run the console consumer commands on terminal to view the output of the program, using the below command:

- **To read contents of ItemTopic**

```
./bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning \
--zookeeper localhost:2181 --property print.key=true
```

- **To read contents of UserTopic**

```
./bin/kafka-console-consumer.sh --topic UserTopic --from-beginning \
--zookeeper localhost:2181 \
--property print.key=true
```



```
[acadgild@localhost kafka_2.12-0.10.1.1]$ ./bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning \
> --zookeeper localhost:2181 \
> --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"item_id":"101"} {"user_id":"U101"}
{"item_id":"101"} {"user_id":"U106"}
{"item_id":"102"} {"user_id":"U110"}
{"item_id":"102"} {"user_id":"U101"}
{"item_id":"104"} {"user_id":"U102"}
{"item_id":"107"} {"user_id":"U104"}
^CProcessed a total of 6 messages
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost kafka_2.12-0.10.1.1]$ ./bin/kafka-console-consumer.sh --topic UserTopic --from-beginning \
> --zookeeper localhost:2181 \
> --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"name":"John"} {"exp":16}
{"name":"Mark"} {"exp":18}
{"name":"Cylin"} {"exp":15}
{"name":"Prod"} {"exp":14}
{"name":"Abhay"} {"exp":17}
{"name":"Misano"} {"exp":19}
```

output from console consumer. This is showing contents of "ItemTopic"

Output from console consumer. This is showing contents of "UserTopic"

Task 2: Modify the previous program MyKafkaProducer.java and create a new Java program KafkaProducerWithAck.java. This should perform the same task as of KafkaProducer.java with some modification. When passing any data to a topic, it should wait for acknowledgement. After acknowledgement is received from the broker, it should print the key and value which has been written to a specified topic. The application should attempt for 3 retries before giving any exception.

Pass dataset_producer.txt as the input file and -as delimiter.

Program to perform this task is as below:

- Imports required for the program is given below:

```
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import java.io.BufferedReader;
import java.io.FileReader;
import java.io.IOException;
import java.util.Properties;
import java.util.concurrent.ExecutionException;
```

- This is the class called “KafkaProducerWithAck” which takes two arguments (Input File name and delimiter) in the command line.

```
public class KafkaProducerWithAck {

    public static void main(String[] args) throws IOException, InterruptedException,
    ExecutionException{

        if (args.length != 2) {

            System.out.println("Please provide appropriate command line arguments");

            System.exit(-1);

        }
    }
}
```

- We configure the properties for KafkaProducer:
 - We create a new instance of Properties called *props*.
 - Using this instance we add properties to kafkaProducer like, bootstrapserver/meta-data-brokerlist, key and value serializers,acks and retries.
 - Acks “all”**- this means that the producer will receive a success response from the broker once all in-sync replicas received the message.
 - Retries 3**- When the producer receives an error message from the server, the error could be transient (e.g., a lack of leader for a partition). In this case, the value of the retries parameter will control how many times the producer will retry sending the message before giving up and notifying the client of an issue.

```

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("acks", "all");
props.put("retries", 3);
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

```

- We then instantiate the KafkaProducer class called producer, we have mentioned string in <> because both key and value are String.
- We add the properties instance (props) to KafkaProducer instance.
- We also instantiate ProducerRecord as producerRecord

```
KafkaProducer<String, String> producer = new KafkaProducer<>(props);
```

```
ProducerRecord<String, String> producerRecord = null;
```

- Now we take the data provided in the command line i.e. file name and delimiter and save them in the array of string variables called filename and delimiter

```
String fileName = args[0];
```

```
String delimiter = args[1];
```

- We read the contents of the input file, and save their contents arrays in different variables:
 - We save the topic name i.e. first part of array(0th index elements) in String variable topic and similarly we save key and value variables too.

```
try(BufferedReader br = new BufferedReader(new FileReader(fileName))) {
```

```
    for(String line; (line = br.readLine()) != null; ) {
```

```
        String[] tempArray = line.split(delimiter);
```

```
        String topic = tempArray[0];
```

```
        String key = tempArray[1];
```

```
        String value = tempArray[2];
```

- Now, we pass the variables topic, key and value to producer record.
- We also print appropriate message which shows the topics, key and value contents.
- We finally, close the producer

```

producerRecord = new ProducerRecord<String, String>(topic, key, value);

producer.send(producerRecord).get();

System.out.printf("Record sent to topic:%s and acknowledged as well. Key:%s,
Value:%s\n", topic, key, value);

} } producer.close(); } }

```

- We start zookeeper and kafka server using the below commands:

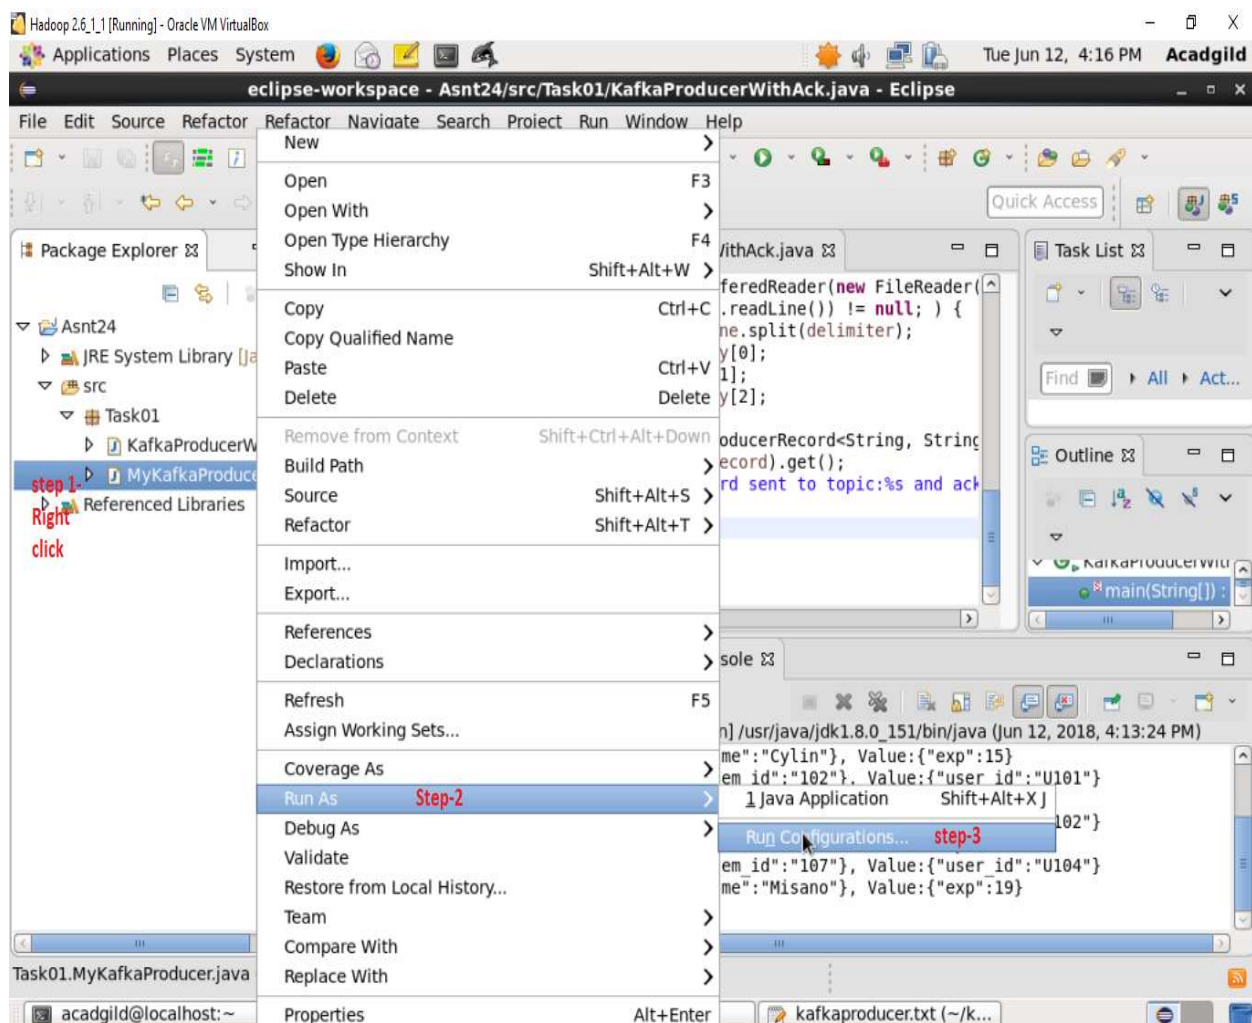
- ## starting zookeeper

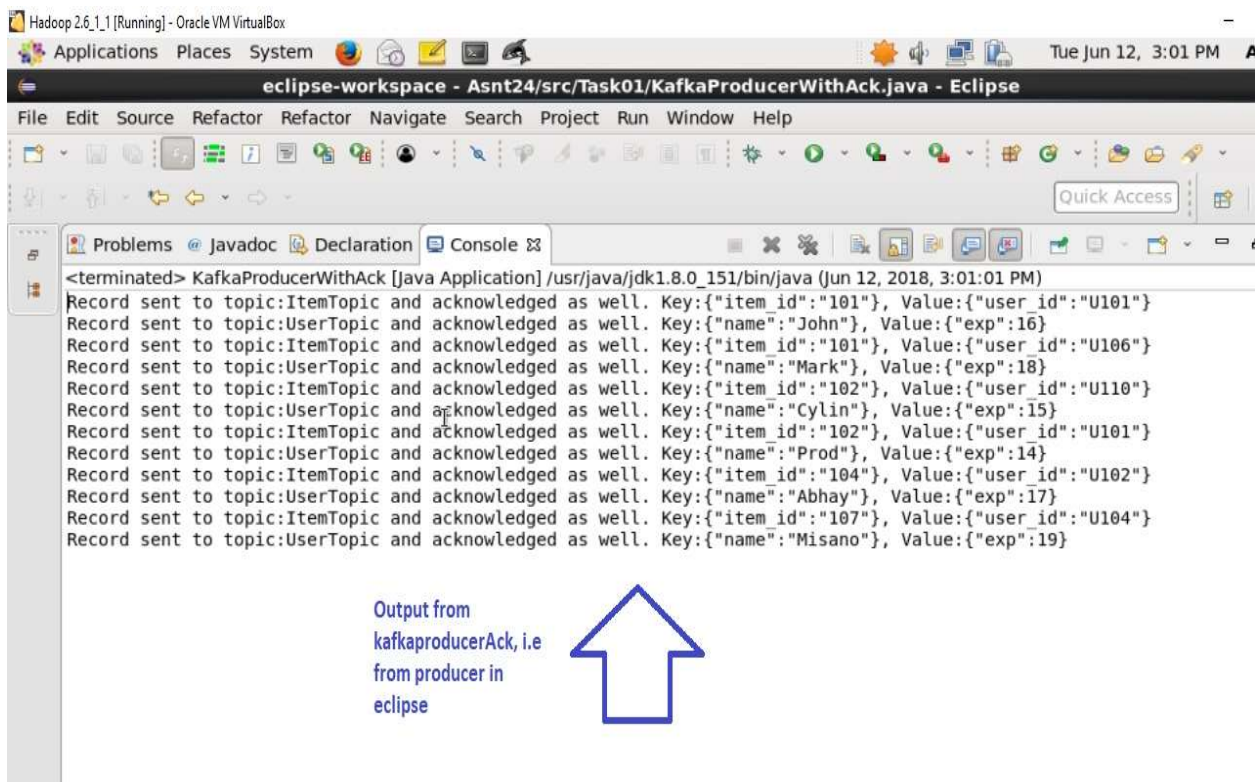
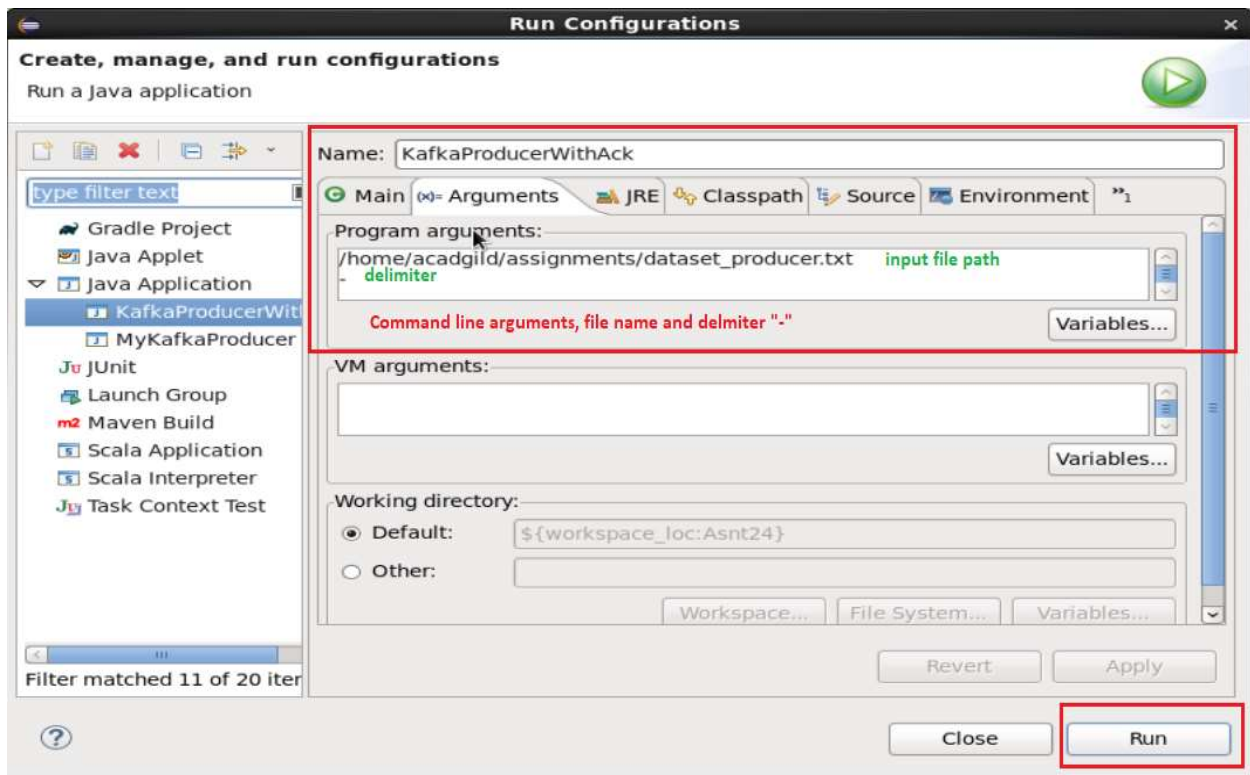
bin/zookeeper-server-start.sh config/zookeeper.properties

- ## starting kafka server

bin/kafka-server-start.sh config/server.properties

- Now, we run this program in eclipse, by giving the arguments in “Run Configurations” as shown below:



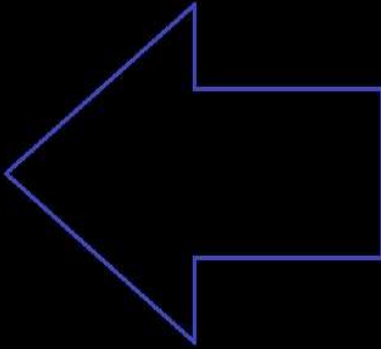


- Now, we run the console consumer commands on terminal to view the output of the program, using the below command:
 - **To read contents of ItemTopic**

`./bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning \`
`--zookeeper localhost:2181 --property print.key=true`

```
[acadgild@localhost kafka 2.12-0.10.1.1]$ ./bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning \
> --zookeeper localhost:2181 \
> --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new
consumer by passing [bootstrap-server] instead of [zookeeper].
{"item_id":"101"}      {"user_id":"U101"}
{"item_id":"101"}      {"user_id":"U106"}
{"item_id":"102"}      {"user_id":"U110"}
{"item_id":"102"}      {"user_id":"U101"}
{"item_id":"104"}      {"user_id":"U102"}
{"item_id":"107"}      {"user_id":"U104"}
{"item_id":"101"}      {"user_id":"U101"}
{"item_id":"101"}      {"user_id":"U106"}
{"item_id":"102"}      {"user_id":"U110"}
{"item_id":"102"}      {"user_id":"U101"}
{"item_id":"104"}      {"user_id":"U102"}
{"item_id":"107"}      {"user_id":"U104"}
{"item_id":"101"}      {"user_id":"U101"}
{"item_id":"101"}      {"user_id":"U106"}
{"item_id":"102"}      {"user_id":"U110"}
{"item_id":"102"}      {"user_id":"U101"}
{"item_id":"104"}      {"user_id":"U102"}
{"item_id":"107"}      {"user_id":"U104"}
{"item_id":"101"}      {"user_id":"U101"}
{"item_id":"101"}      {"user_id":"U106"}
{"item_id":"102"}      {"user_id":"U110"}
{"item_id":"102"}      {"user_id":"U101"}
{"item_id":"104"}      {"user_id":"U102"}
{"item_id":"107"}      {"user_id":"U104"}
{"item_id":"101"}      {"user_id":"U101"}
{"item_id":"101"}      {"user_id":"U106"}
{"item_id":"102"}      {"user_id":"U110"}
{"item_id":"102"}      {"user_id":"U101"}
{"item_id":"104"}      {"user_id":"U102"}
{"item_id":"107"}      {"user_id":"U104"}
```

Output of console consumer for "ItemTopic"



- To read contents of UserTopic

```
./bin/kafka-console-consumer.sh --topic UserTopic --from-beginning \  
--zookeeper localhost:2181 \  
--property print.key=true
```

```
[acadgild@localhost kafka_2.12-0.10.1.1]$ ./bin/kafka-console-consumer.sh --topic UserTopic --from-beginning \  
> --zookeeper localhost:2181 \  
> --property print.key=true  
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new c  
onsumer by passing [bootstrap-server] instead of [zookeeper].  
{ "name": "John" } { "exp": 16 }  
{ "name": "Mark" } { "exp": 18 }  
{ "name": "Cylin" } { "exp": 15 }  
{ "name": "Prod" } { "exp": 14 }  
{ "name": "Abhay" } { "exp": 17 }  
{ "name": "Misano" } { "exp": 19 }  
{ "name": "John" } { "exp": 16 }  
{ "name": "Mark" } { "exp": 18 }  
{ "name": "Cylin" } { "exp": 15 }  
{ "name": "Prod" } { "exp": 14 }  
{ "name": "Abhay" } { "exp": 17 }  
{ "name": "Misano" } { "exp": 19 }  
{ "name": "John" } { "exp": 16 }  
{ "name": "Mark" } { "exp": 18 }  
{ "name": "Cylin" } { "exp": 15 }  
{ "name": "Prod" } { "exp": 14 }  
{ "name": "Abhay" } { "exp": 17 }  
{ "name": "Misano" } { "exp": 19 }  
{ "name": "John" } { "exp": 16 }  
{ "name": "Mark" } { "exp": 18 }  
{ "name": "Cylin" } { "exp": 15 }  
{ "name": "Prod" } { "exp": 14 }  
{ "name": "Abhay" } { "exp": 17 }  
{ "name": "Misano" } { "exp": 19 }  
{ "name": "John" } { "exp": 16 }  
{ "name": "Mark" } { "exp": 18 }  
{ "name": "Cylin" } { "exp": 15 }  
{ "name": "Prod" } { "exp": 14 }  
{ "name": "Abhay" } { "exp": 17 }  
{ "name": "Misano" } { "exp": 19 }
```

Output of console consumer for the "UserTopic"

